



RAG basics

JUNE 2025

There are no magic! It is simple REST API!

Before we start:


- **Raise your hand and ask questions if you have any**
- **It is better to ask questions when you have**
- **Also, type them in chat**
- **We will need DIAL API key for this session**

Agenda:

- **Presentation:**
 - **About RAG Concept**
 - **Real life sample**
 - **Application architecture and main components**
- **Workshop:**
 - **Create MicrowaveRAG application**

RAG Concept

RAG = Retrieval-Augmented Generation



Retrieval-augmented generation (RAG) is a technique that enables Gen AI models to retrieve and incorporate new information. It modifies interactions with a LLM so that the model responds to user queries with reference to a specified set of documents, using this information to supplement information from its pre-existing training data. This allows LLMs to use domain-specific and/or updated information. Use cases include providing chatbot access to internal company data or generating responses based on authoritative sources.




R A G concept:



Retrieval

- The system searches through an external knowledge base (documents, databases, webpages, or vector stores) to find information relevant to the user query.
- Often, this is done using vector embeddings (semantic search) to find relevant documents based on similarity measures.


R A G concept:



The retrieved information is then used to extend or "augment" the context provided to the language model.

«User input + Retrieved data»

R A G concept:



**The LLM generates response based on the provided information
(user input + retrieved data)**

Retrieval-augmented generation (RAG) is a technique that helps us:



Provide the most relevant context data based on user request

Reduces hallucinations

Enables up-to-date knowledge

Enables domain specialization

Reduces context window usage

Reduces costs (not always*)

Sample

User request

Wash all
the red
clothes



User request

Wash all
the red
clothes

Retrieval

Got it.
Now I'll
retrieve all
the red
clothes
from this
bucket



User request

Wash all
the red
clothes

Retrieval

Got it.
Now I'll
retrieve all
the red
clothes
from this
bucket

Here are
all the red
clothes
that I was
able to find



User request

Wash all
the red
clothes

Retrieval

Got it.
Now I'll
retrieve all
the red
clothes
from this
bucket

Here are
all the red
clothes
that I was
able to find

Augmentation

'Wash all the
red clothes'
+ the clothes

User request

Wash all
the red
clothes

Retrieval

Got it.
Now I'll
retrieve all
the red
clothes
from this
bucket

Here are
all the red
clothes
that I was
able to find

Augmentation

'Wash all the
red clothes'
+ the clothes

Generation

I've made the laundry with all
the provided red stuff. There
were a couple of rags and a pair
of shoes. Any other
instructions?

Wash all the
red clothes


User request



Retrieval



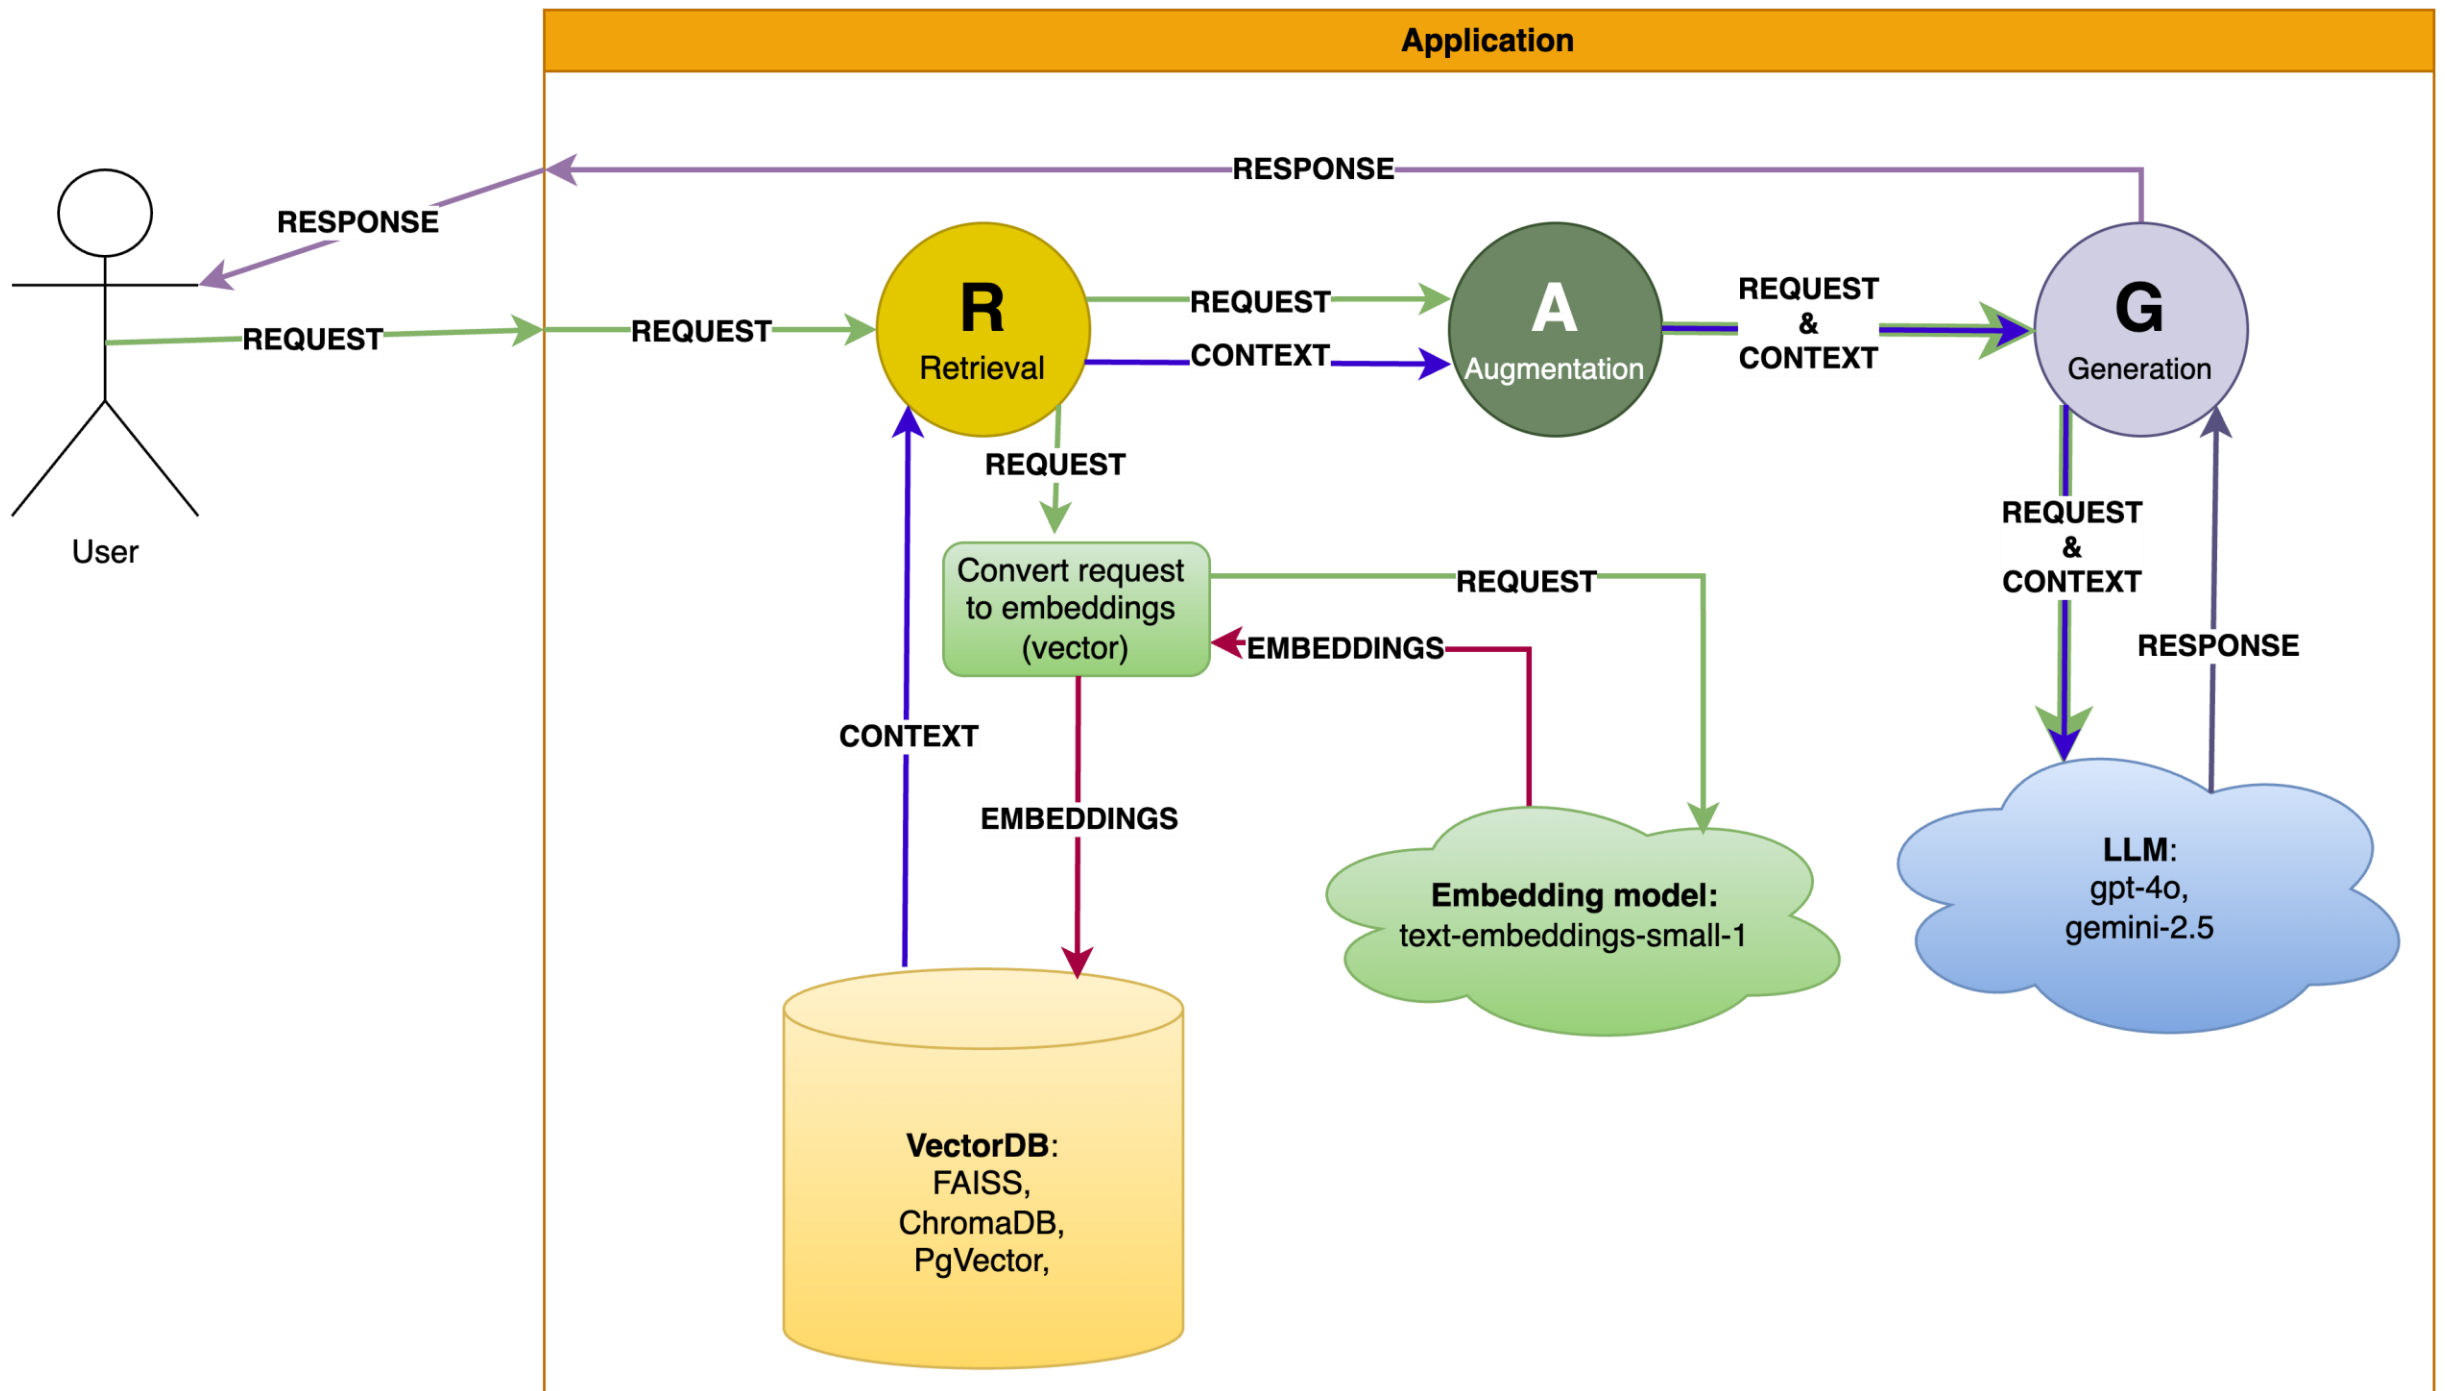
Augmentation



Generation



Application



Client that is working with Embedding model. Via this client you will convert text chunks and user requests into embeddings.

Embedding models:

- text-embedding-3-large-1
- text-embedding-3-smal-1

Client to work with LLM:

- gpt-4o
- gemini-2.5
- sonnet-3.7

```
AzureOpenAIEmbeddings(  
    deployment='text-embedding-3-large-1',  
    azure_endpoint=DIAL_URL,  
    api_key=SecretStr(API_KEY),  
)
```

```
AzureChatOpenAI(  
    temperature=0.0,  
    azure_deployment='gpt-4o-2024-08-06',  
    azure_endpoint=DIAL_URL,  
    api_key=SecretStr(API_KEY),  
    api_version="2024-05-01-preview"  
)
```

Via this API you will be able to load locally saved indexed FAISS Vector DB. This is needed to avoid creation new DB each time when Application is started

```
FAISS.load_local(  
    folder_path='microwave_faiss_index',  
    embeddings=self.embeddings,  
    allow_dangerous_deserialization=True,  
)
```

Will convert Chunks (documents) into embeddings and create FAISS Vector DB with them. Then you can locally save it.

```
FAISS.from_documents(chunks, self.embeddings)
```

Will convert Chunks (documents) into embeddings and create FAISS Vector DB with them. Then you can locally save it.

```
RecursiveCharacterTextSplitter(  
    chunk_size=300,  
    chunk_overlap=50,  
    separators=["\n\n", "\n", "."]  
)
```

Via this API you can make similarity search relevant context to user request (query).

- `k` – limit of results that we expect to get
- `score_threshold` – filter with similarity score. Range 0.0-1.0

```
vectorstore.similarity_search_with_relevance_scores(  
    query,  
    k=k,  
    score_threshold=score  
)
```




Join us:



Subscribe to WeAreCommuntiy

<https://wearecommunity.io/communities/dial>



Keep in touch with our latest updates.
Here you find webinars, workshops and
articles about DIALX features and products.

Subscribe to YouTube

<https://www.youtube.com/@TeamDIALX>



Here we publish videos about our newest
products and features.

Join our Discord community

<https://discord.gg/jvTCQv4E4q>



🌟 AI DIALX Community 🌟 is the place where
you can find help with your questions about
DIALX, direct communication with DIALX team
and contributors.



DIALX
COMMUNITY
POWERED BY <epam>

Thank you!